

Estimation of Human-perceived Difficulty of Chess Positions

Zala Urh

zu0476@student.uni-lj.si
Faculty of Computer and
Information Science,
University of Ljubljana
Večna pot 113
SI-1000 Ljubljana, Slovenia

Matej Belšak

mb4390@student.uni-lj.si
Faculty of Computer and
Information Science,
University of Ljubljana
Večna pot 113
SI-1000 Ljubljana, Slovenia

Luka Bajić

lb4129@student.uni-lj.si
Faculty of Computer and
Information Science,
University of Ljubljana
Večna pot 113
SI-1000 Ljubljana, Slovenia

ABSTRACT

This paper overviews various methods for predicting the difficulty of a given chess position from the viewpoint of human players. Feature engineering is implemented via limited-depth search with a chess engine, with the emphasis on interpretability. Modeling and evaluation are split into classification, regression and ranking.

KEYWORDS

chess, classification, regression, ranking, feature engineering

1 INTRODUCTION

In this paper we compare various methods for determining the difficulty of a given chess position for human players, using chess engine evaluations and other attributes of the position. This problem can be defined in the following three ways:

- classification: the model classifies positions into pre-defined classes (e.g. easy/medium/hard), or binned ELO ratings.
- regression: the model directly predicts an estimated ELO rating of the position.
- ranking: the model orders a set of positions from easiest to hardest, or ranks a single position in relation to all positions in the training set.

Full implementation is available at ¹.

2 RELATED WORK

Recent approaches for estimating the difficulty of a chess position, e.g. [4] are largely based on the deep learning paradigm and the input of these models consists of some encoding of the chess position, but no information that would provide interpretability for human players. In this paper we therefore focus on finding features that achieve good classification score, but also provide a meaningful interpretation regarding the reasons why a position is considered difficult or easy for human players.

Our approaches are based on insights provided by [1], [2], [5]. These studies developed numerous attributes that can be obtained by performing a chess engine search of certain depth, starting in the initial position we wish to evaluate. The search is constrained in such a way that only moves which human players would typically assess as meaningful are explored, which aims to provide an estimate of how complex the search tree around the position is.

3 METHODOLOGY

To train and evaluate our models, we use a subset of the Lichess puzzle dataset ² containing 1000 puzzles with ratings ranging between 399 and 3092. For classification we create classes "easy", "medium" and "hard" by splitting this interval equidistantly. The dataset is split into training set and test set with the ratio 80:20. The data is standardized, meaning we transform it in such a way that it follows a standard normal distribution with zero mean and unit variance. This ensures that each attribute is treated as equally important by the models.

3.1 Evaluation Metrics

For classification problem we use the standard metrics: classification accuracy, F1 score and AUC. For regression we use mean absolute error (MAE) because unlike mean squared error, it has a more intuitive interpretation on this dataset. We also report which percentage of the rating interval this MAE covers. For ranking we use Normalized Discounted Cumulative Gain (NDCG) and Kendall τ to determine the accuracy of the model's ranking against the ground truth, which is simply a ranking assigned by ordering the ratings in the dataset.

3.2 Attributes

Other than a list of themes for each puzzle, the original dataset contains no attributes that would allow modeling of perceived difficulty for human players, so we expand it with the following attributes:

- Stockfish parameters for each depth between 1-3 from starting position. These parameters include centipawn evaluations, number of searched nodes, WDL (win/draw/loss) and selective depth.
- material count and piece counts for each side.
- attributes from [5], computed for depth=3, excluding "Seemingly Good" due to ambiguous definition.
- binary attributes for solving success by Stockfish engine with different playing strengths.

A more detailed explanation of the attributes is given in Appendix 1. Interpretability of the attributes is discussed in Appendix 3.

¹<https://github.com/bajicluka01/chess-puzzle-evaluation/tree/main>

²<https://database.lichess.org/#puzzles>

4 RESULTS

4.1 Classification

Table 1 displays the performance of the standard machine learning models on a dataset with 1000 samples and all available attributes. Class distribution is 39.1% "easy", 38.9% "medium" and 22.0% "hard".

XGBoost uses 100 estimators and learning rate 0.1. SVM uses regularization parameter $C = 1$ and balanced class weight. Random Forest uses 100 estimators. kNN uses 10 nearest neighbors. MLP uses $layers = [2048, 1024, 512, 256, 128]$, ReLU activation function, ADAM solver, regularization $\alpha = 0.001$ and batch size 512. Default values are used for all remaining parameters. Ensemble uses a soft voting scheme on the outputs of all other models, meaning that probabilistic predictions are aggregated and finally converted into a class value. Misclassifications are discussed in Appendix 2.

Table 1: Classification scores on 1k dataset.

model	CA	F1	AUC
XGBoost	0.575	0.574	0.756
kNN	0.635	0.621	0.773
Random Forest	0.620	0.616	0.772
Gradient Boosting	0.575	0.563	0.758
MLP	0.630	0.621	0.770
SGD	0.550	0.551	0.693
SVM	0.595	0.574	0.756
Ensemble	0.615	0.612	0.779

In Table 2 we show the results of an ablation study aimed to determine the utility of each subset of attributes. We use the best performing model in terms of CA from the previous experiment: kNN with $k = 10$. Baseline refers to simple attributes, such as centipawn evaluations and number of searched nodes for each depth. These attributes are used in each experiment. Study refers to attributes from [5].

Table 2: Ablation study on 1k dataset.

model	CA	F1	AUC
baseline	0.565	0.561	0.671
counts	0.555	0.545	0.688
study	0.565	0.548	0.716
Stockfish	0.565	0.557	0.752
themes	0.565	0.554	0.704
counts, study	0.530	0.508	0.723
counts, Stockfish	0.575	0.563	0.716
counts, themes	0.565	0.553	0.713
study, Stockfish	0.580	0.561	0.766
study, themes	0.58	0.570	0.717
all attributes	0.635	0.621	0.773

This confirms the usefulness of attributes from [5] and also confirms the usefulness of themes. The contribution of solving success

is negligible, possibly due to an inherent flaw in the implementation of Stockfish playing strength: weaker play is simulated by limiting the depth of the search tree and imposing a time constraint on the search, however, this might not accurately simulate the puzzle-solving ability of a player with a certain ELO, rendering the attributes useless.

4.2 Regression

Regression results are shown in Table 3. XGBoost uses 100 estimators and learning rate 0.1. MLP uses $layers = [256, 256, 128, 128, 128, 64]$, ReLU activation function, adaptive learning rate, ADAM solver, regularization parameter $\alpha = 0.001$ and batch size 256. Default values are used for all remaining parameters. Ensemble model performs unweighted averaging on the outputs of all other models.

Table 3: Regression scores on 1k dataset.

model	MAE	% of interval
XGBoost	329.836	12.2
Linear regression	321.600	11.9
MLP	317.419	11.8
Bayes Ridge	314.587	11.7
Ensemble	310.385	11.5

4.3 Ranking

Ranking results are shown in Table 4. XGBRanker uses 100 estimators and learning rate 0.1, CatBoostRanker uses 1000 iterations and learning rate 0.1. Default values are used for all remaining parameters.

Table 4: Ranker scores on 1k dataset.

model	NDCG	τ
XGBRanker	0.84	0.41
LGBMRanker	0.87	0.43
CatBoostRanker	0.85	0.48

Kendall τ is a metric defined on $[-1, 1]$, so these scores indicate that the rankers do capture at least some notion of difficulty from the attributes.

5 CONCLUSIONS

Future work would consist of checking if these approaches are suitable also for strategic/nondescript positions, rather than just puzzles. Here, the main challenge is in assigning the ground truth ELO to these positions, since consensus among human players is more difficult to achieve.

Additionally, Maia chess engine [3] [6] could be used as an alternative to Stockfish, as it is designed to approximate human play more closely. However, we ran a few experiments on known positions, and again this similarity to human-like play does not directly translate to human-like puzzle solving, so additional model training may be required, which is outside the scope of this paper.

REFERENCES

- [1] Ivan Bratko, Dayana Hristova, and Matej Guid. 2016. *Search Versus Knowledge in Human Problem Solving: A Case Study in Chess*. Vol. 27. 569–583. https://doi.org/10.1007/978-3-319-38983-7_32
- [2] Ivan Bratko, Dayana Hristova, and Matej Guid. 2021. Predicting Problem Difficulty in Chess. In *Human-Like Machine Intelligence*, Stephen Muggleton and Nicholas Chater (Eds.). Oxford University Press. <https://doi.org/10.1093/oso/9780198862536.003.0024> arXiv:<https://academic.oup.com/book/0/chapter/350718010/chapter-pdf/43431216/oso-9780198862536-chapter-24.pdf>
- [3] Reid McIlroy-Young, Siddhartha Sen, Jon Kleinberg, and Ashton Anderson. 2020. Aligning Superhuman AI with Human Behavior: Chess as a Model System. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '20)*. ACM, 1677–1687. <https://doi.org/10.1145/3394486.3403219>
- [4] Szymon Milosz and Pawel Kapusta. 2024. Predicting Chess Puzzle Difficulty with Transformers. arXiv:2410.11078 [cs.LG] <https://arxiv.org/abs/2410.11078>
- [5] Simon Stoiljkovic, Ivan Bratko, and Matej Guid. 2015. A Computational Model for Estimating the Difficulty of Chess Problems. <https://api.semanticscholar.org/CorpusID:204857616>
- [6] Zhenwei Tang, Difan Jiao, Reid McIlroy-Young, Jon Kleinberg, Siddhartha Sen, and Ashton Anderson. 2024. Maia-2: A Unified Model for Human-AI Alignment in Chess. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*. <https://openreview.net/forum?id=XWlkhRn14K>

APPENDIX 1: ATTRIBUTES

Detailed list of attributes used in the ablation study (all attributes combined are used for all other experiments):

- baseline:
 - to_move: which side moves first
 - cp_eval: Stockfish’s centipawn evaluation of the initial position
 - moveXcp: Stockfish’s centipawn evaluation of the position after X moves have been made (only best move is considered at each turn).
 - nodes:
 - selective depth: number of half-moves calculated by the engine in the most forcing variation.
 - WDL: win/draw/loss probability reported by the engine
- counts: total material count for each player, counts for each piece for each player (e.g. number of white queens on the board).
- study:
 - number of meaningful moves at level L: an attempt to count the number of moves human players would consider meaningful in the given position.
 - branching at level L: branching factor that takes into account only meaningful moves.
 - average branching: an average of branching factors over all layers.
 - narrow solution at level L: counts the number of "forced" moves (that is, positions in which only one move leads to a good position).
 - distance at level L: length between starting and target square for each move (that is, actual number of squares between the two squares).
 - pieces at level L: number of different pieces that move at level L.
 - all pieces involved: number of different pieces that move in the meaningful subtree.

- winning without checkmate: number of moves evaluated as winning that do not directly lead to a checkmate.
- possible moves at level L: total number of legal moves at level L.
- all possible moves: total number of legal moves at all levels.
- all narrow solutions: sum of narrow solutions over all levels.
- tree size: number of nodes in the meaningful search tree.
- move ratio at level L: ratio between meaningful moves and possible moves.
- sum of distances over all levels.
- average distance over all levels.
- Stockfish: engines with various strength settings (between 1320 and 3000) are used to evaluate the position. If their solution is aligned with the ground truth solution in the dataset, these attributes are given value 1, otherwise 0.
- themes: 63 theme indicators provided in the original Lichess dataset. These include, for instance, "advancedPawn", "discoveredCheck", "doubleBishopMate", "mateIn4", "trapped-Piece", etc. For each position, multiple themes can be present - for those, the attribute value is 1, otherwise 0.

APPENDIX 2: MISCLASSIFICATIONS




Figure 1: Example of a misclassification between "easy" and "hard": the model classifies this position as easy, although the ground truth label states "hard".

APPENDIX 3: INTERPRETABILITY

Using Figure 2, we can interpret the model’s classification decisions based on a few most important attributes: low selective depth leads to low rating predictions, which makes intuitive sense, because a weaker human player is also likely to find a short forced line.

Solved1320 indicates whether Stockfish engine with strength 1320 was able to solve the given puzzle. If yes, the predicted rating

will be lower, which should also imply that human-perceived difficulty will be lower than if the weak engine was not able to solve the puzzle.

Another obviously helpful attribute is the length of the solution to the puzzle: short puzzles lead to lower rating predictions, which is sensible because lower rated players are on average more likely to successfully solve short puzzles than long ones.




Figure 2: SHAP values for XGBoost Classifier